



ALGORITHM BOOTCAMP

# DEEP LEARNING FOR CHEST X-RAYS

By: Angela Christabel



# TABLE OF CONTENTS

## Presentation Outline

1. Business Background
2. Data Understanding and Exploratory Data Analysis
3. Modelling and evaluation
4. Conclusion and Recommendation





# BUSINESS BACKGROUND

## Background

I am a data scientist working at a health care institution. During this pandemic, it is important, even more so, to be able to successfully diagnose patients who have COVID-19, which is a pneumonia-like disease.



# BUSINESS BACKGROUND

## Business Objective

The objective of this project is to predict whether someone has a pneumonia-like disease or is completely healthy, given their chest x-rays. This will make healthcare workers' jobs easier.



# Project Limitations

## ■ Limited time and resources

Due to this project's limited time and resources, I will only be able to tune my model as far as performing hyperparameter tuning and I won't be able to try various techniques to improve my model.





# Analytic Approach

## ■ Deep Learning Technique

Supervised Learning (classification) to predict whether someone has pneumonia or not.

## ■ Performance Measures

Recall supported by AUC scores to minimize prediction error.





# **Data Understanding and Exploratory Data Analysis**



## Data Collection

# KAGGLE

The dataset I used in this project was downloaded from Kaggle, titled "CoronaHack Chest X-Ray Dataset".


Data shape:

- 5286 rows
- 6 columns





# Data Description



## X\_ray\_image\_name

File names of all X-ray images contained in this dataset.

## Label

Label for each X-ray, whether that X-ray indicates healthy lungs or lungs infected by a pneumonia-like disease.

## Label\_1\_Virus\_category


Label for each X-ray, cause for pneumonia-like disease.

## Label\_2\_Virus\_category

Label for each X-ray, virus/bacteria that caused the pneumonia-like disease.



# Data Description



## Index

Unique index for all X-ray images contained in this dataset.

## Dataset\_type

Specifies whether the certain X-ray image is meant for the training dataset or the test dataset.



# EMPTY VALUES

## All columns

```
Unnamed: 0      0
X_ray_image_name 0
Label           0
Dataset_type     0
Label_2_Virus_category  5217
Label_1_Virus_category  1342
dtype: int64
```

The null values in the Label\_2\_Virus\_category and Label\_1\_Virus\_category columns are probably caused by the fact that the corresponding x-ray images just don't have their causes for pneumonia recorded yet. For further EDA, I filled those empty values with "Unknown".

# VALUES IN COLUMNS



|           |      |
|-----------|------|
| Pneumonia | 3944 |
| Normal    | 1342 |

## Label

These x-rays are labeled either 'Pneumonia' or 'Normal'.

|                |      |
|----------------|------|
| bacteria       | 2535 |
| Virus          | 1407 |
| Unknown        | 1342 |
| Stress-Smoking | 2    |


## Label\_1\_Virus\_category

General causes for the pneumonia-like diseases found in these x-rays.


|               |      |
|---------------|------|
| Unknown       | 5217 |
| COVID-19      | 58   |
| Streptococcus | 5    |
| SARS          | 4    |
| ARDS          | 2    |

## Label\_2\_Virus\_category

Specific causes for the pneumonia-like disease found in these x-rays.



Normal lungs





Lungs infected by  
Pneumonia



# IMAGE HISTOGRAM



Normal lungs





# IMAGE HISTOGRAM


Lungs infected by pneumonia



CORONAHACK



Unique values of Column "Label"



# TARGET CLASS DISTRIBUTION


**Imbalanced**

As we can see from the chart on the left, 'Normal' lungs only make up 25% of the data. But thankfully, the class we care about, 'Pneumonia' is the majority class. Still, we will have to take some measures to handle this imbalance in the data.





# Data Preprocessing




## Methods

After testing using augmentation methods, I concluded that using augmentation methods harmed the model's performance, so the only preprocessing I performed on the data was to resize it to 224 x 224 and rescale the features by dividing each pixel by 255.



ROC Curve

The model without augmentation (Model\_1) is a better classifier on the validation data.




# MODELLING AND EVALUATION





# MODELLING STEPS

Step-by-step approach





```
Model: "sequential"
```

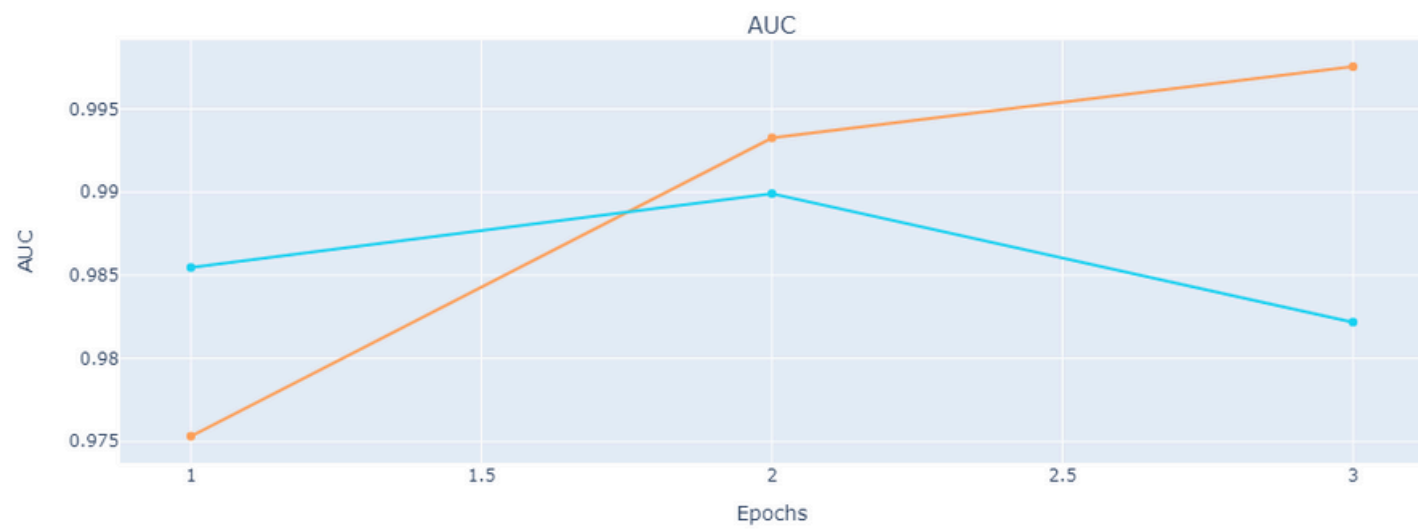
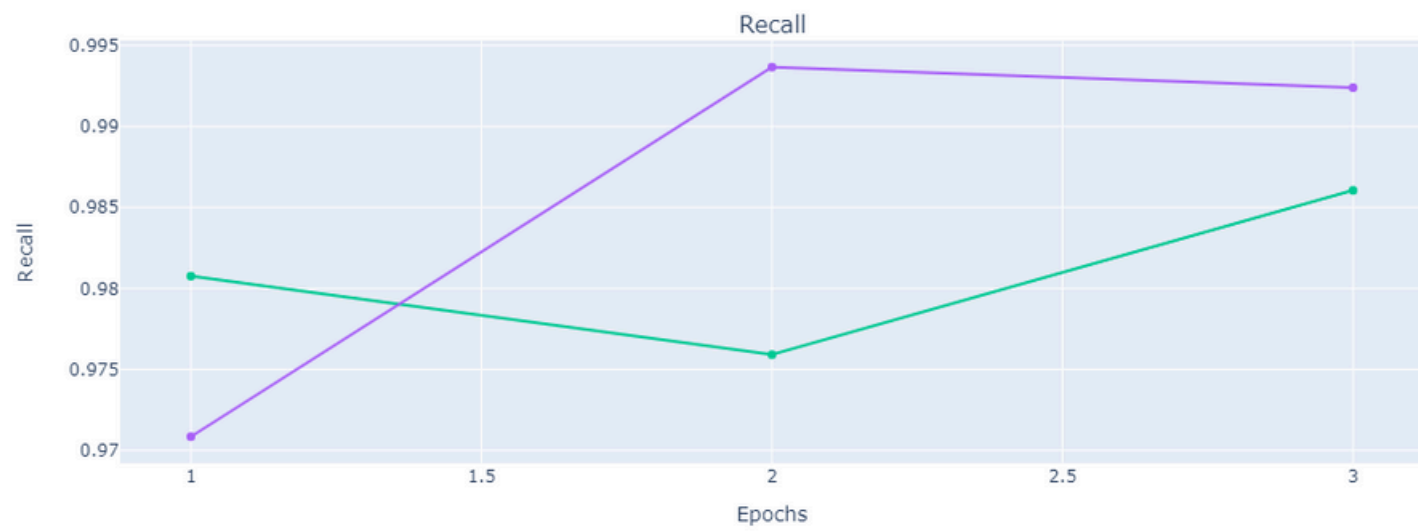
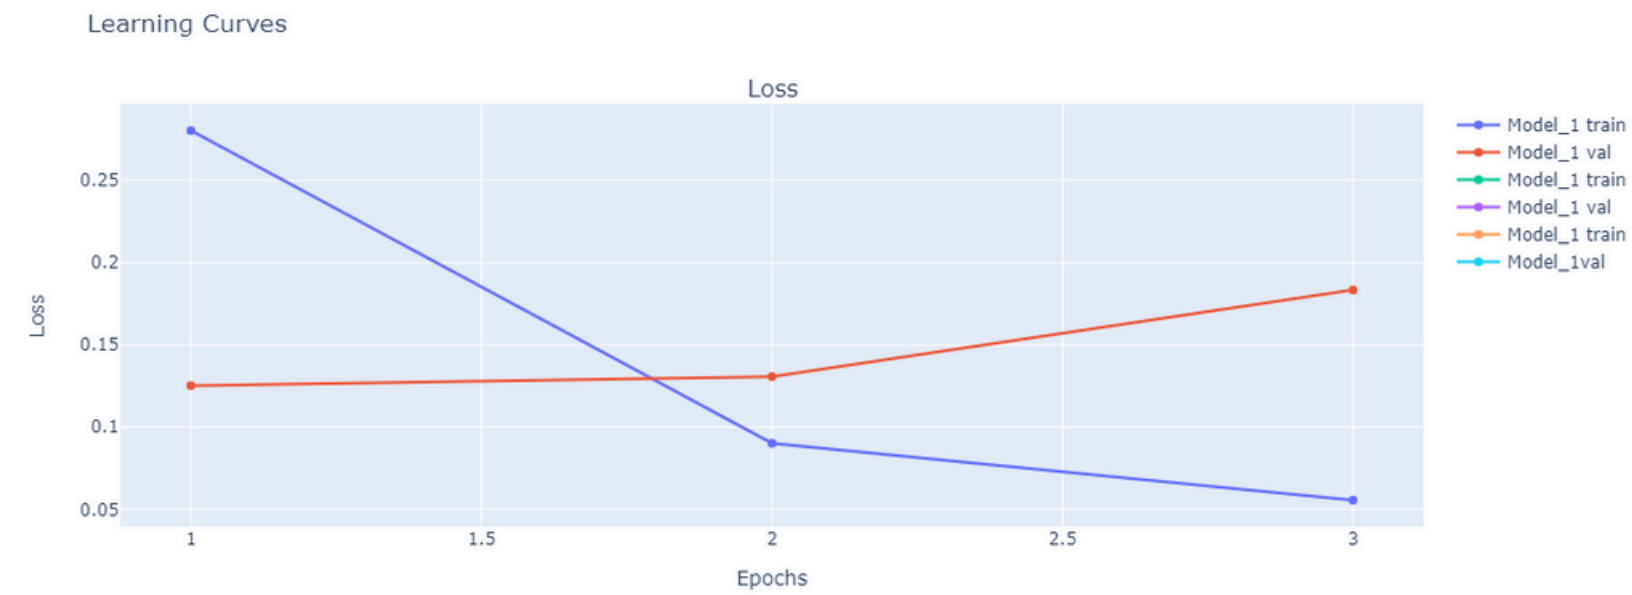
| Layer (type)      | Output Shape         | Param # |
|-------------------|----------------------|---------|
| conv2d (Conv2D)   | (None, 222, 222, 10) | 280     |
| conv2d_1 (Conv2D) | (None, 220, 220, 10) | 910     |
| conv2d_2 (Conv2D) | (None, 218, 218, 10) | 910     |
| flatten (Flatten) | (None, 475240)       | 0       |
| dense (Dense)     | (None, 1)            | 475241  |

```
Total params: 477,341  
Trainable params: 477,341  
Non-trainable params: 0
```

# BASELINE MODEL

## Architecture

I used a 3 layer Convolutional Neural Network as my baseline model.



# BASELINE MODEL

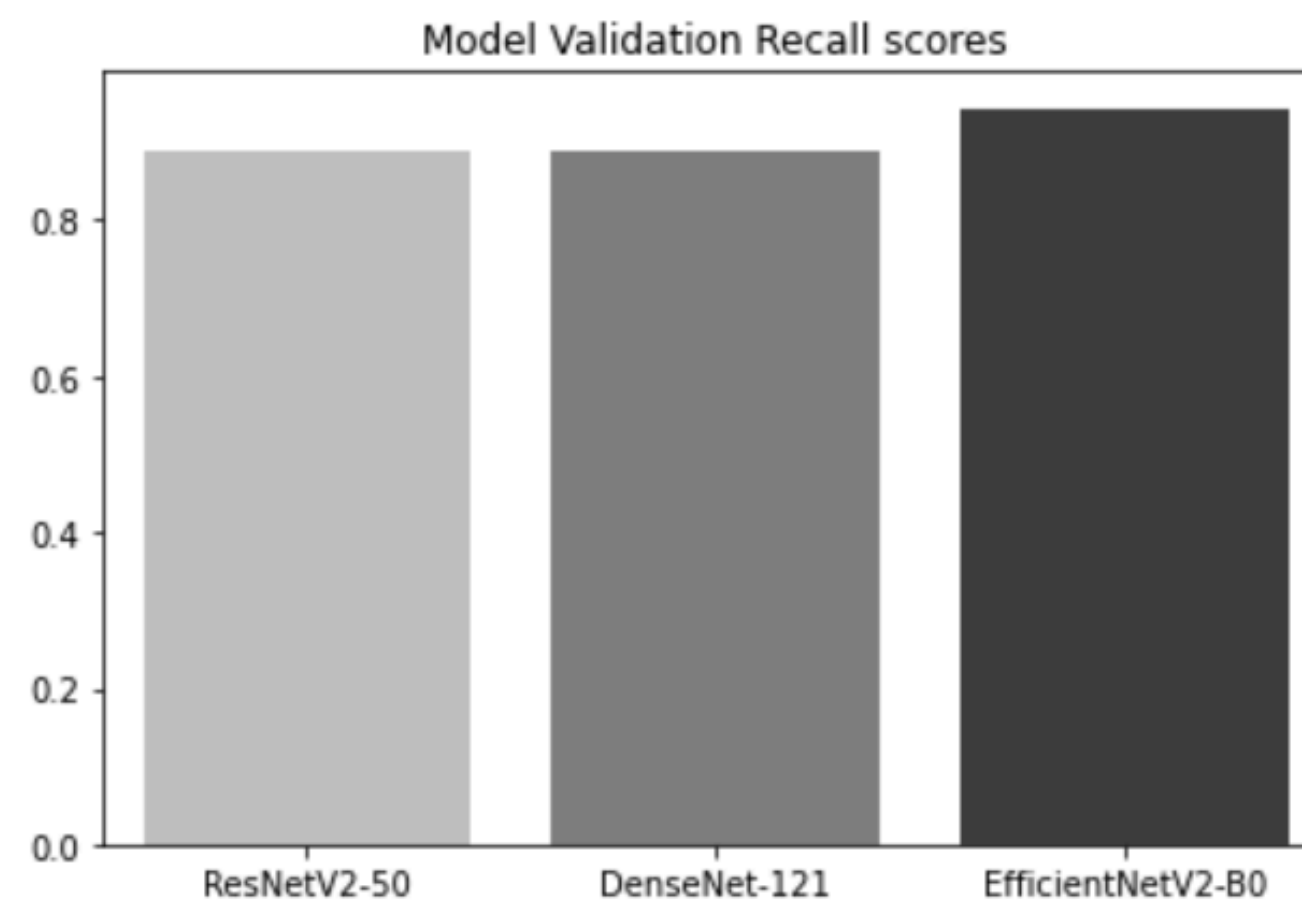
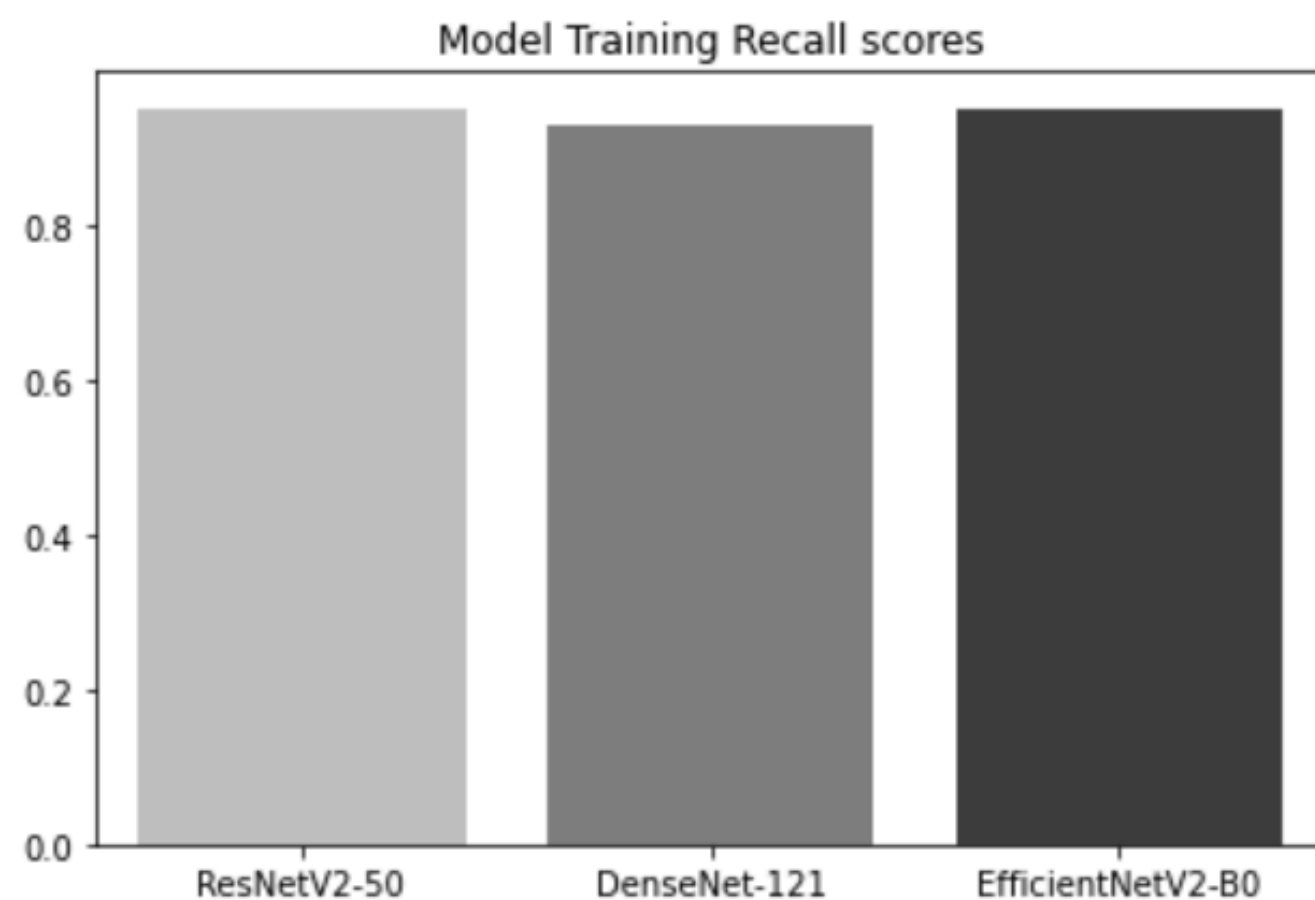
Learning curves



# Pretrained Models' Results

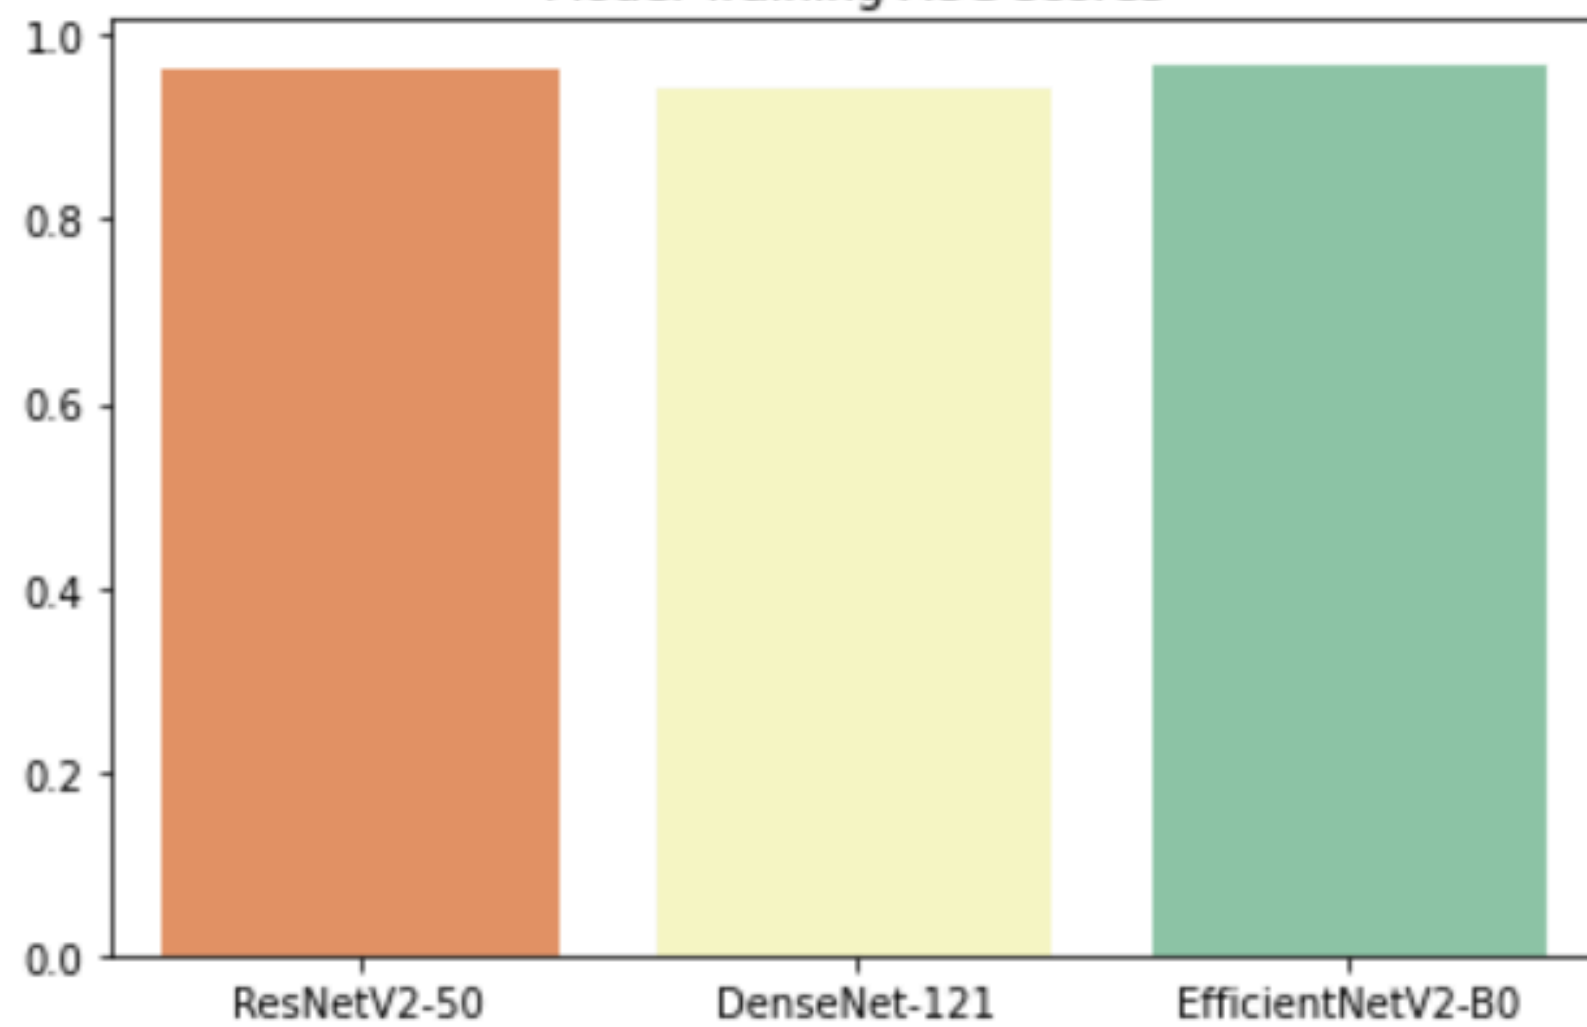


We will focus on the recall score, because it is more important for healthcare workers to be able to lessen the number of wrongly predicted, sick people.

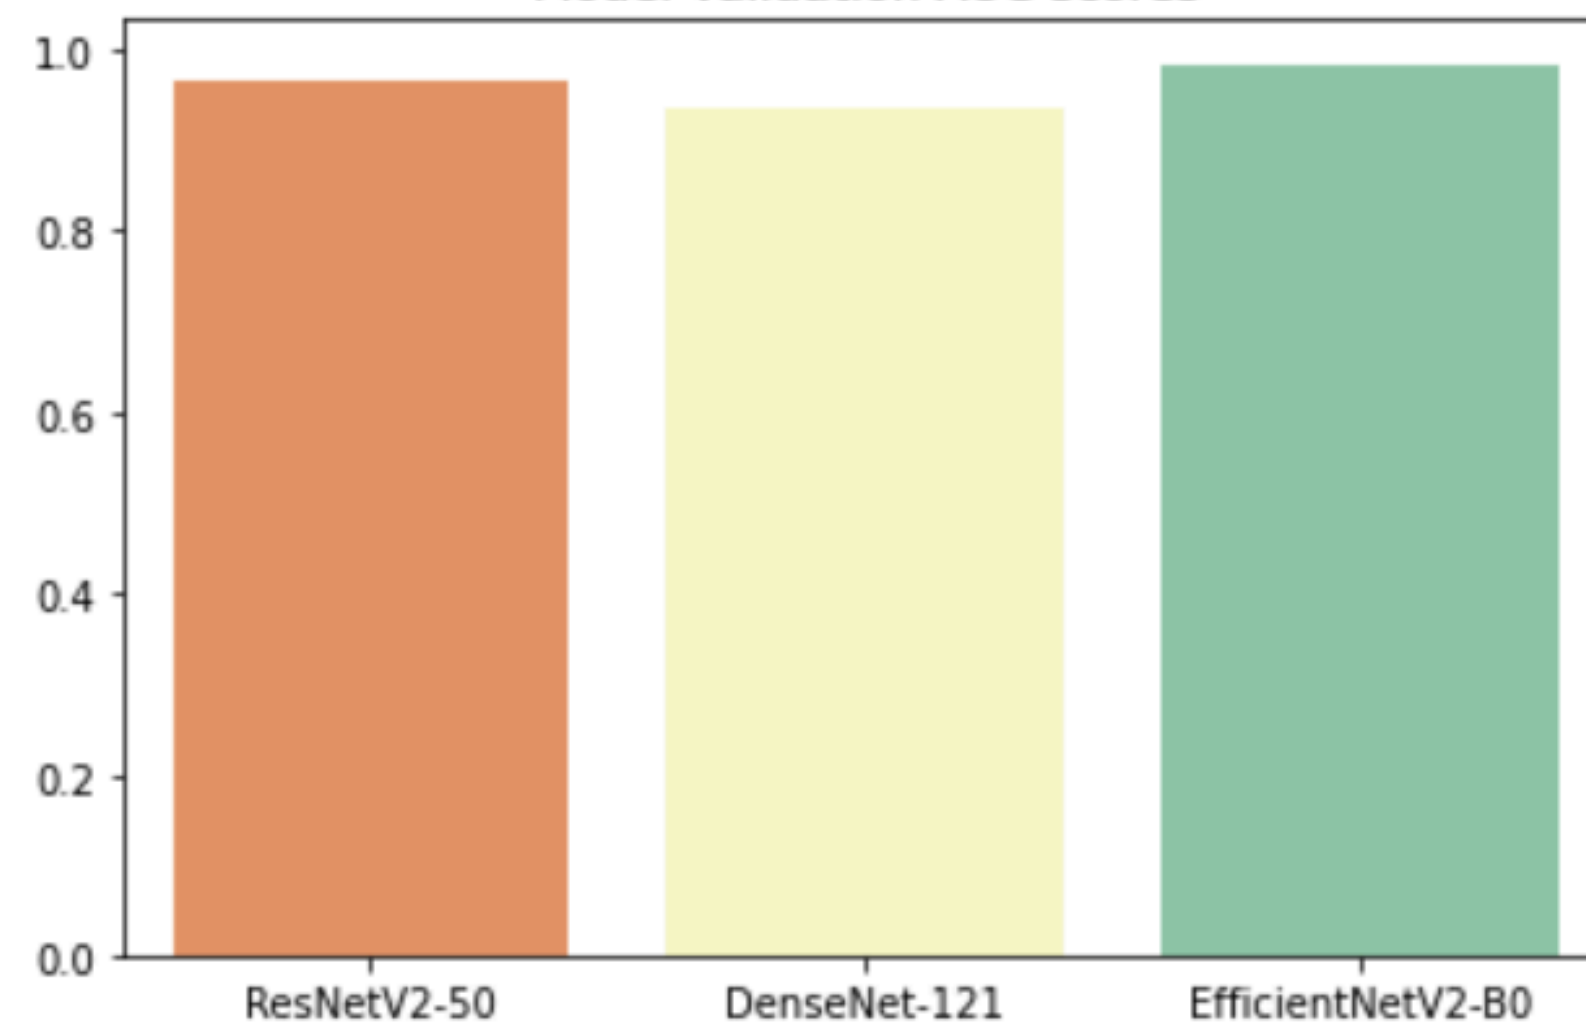


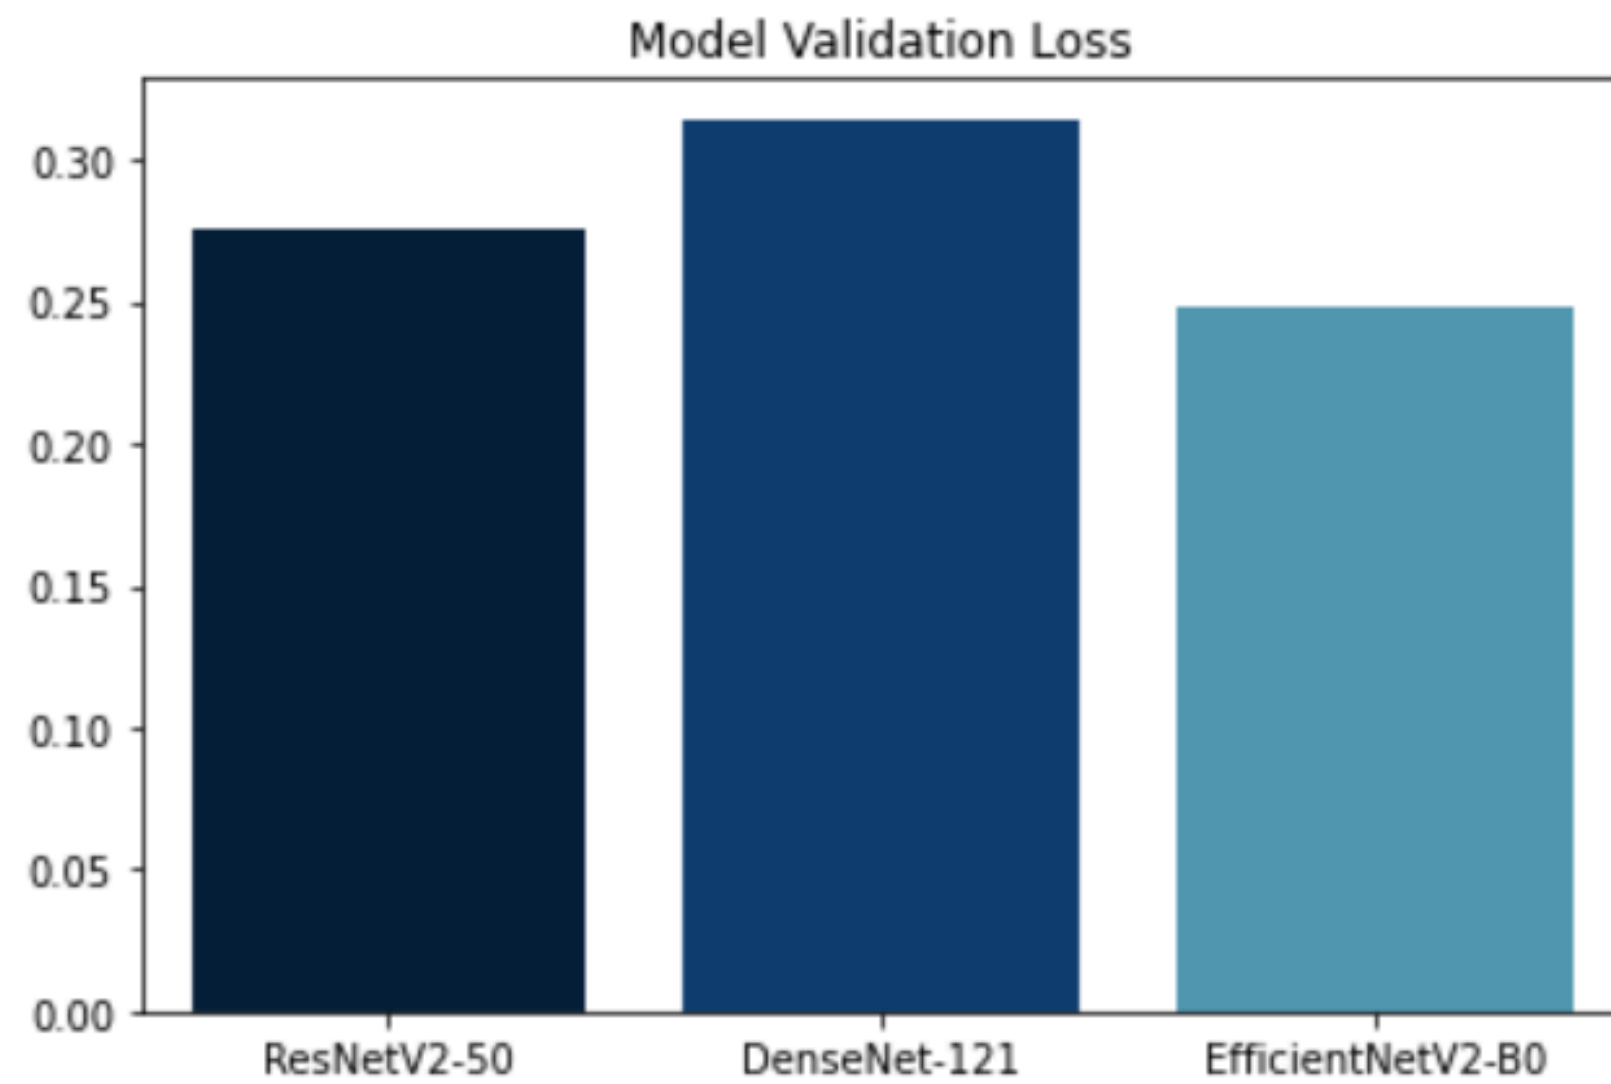
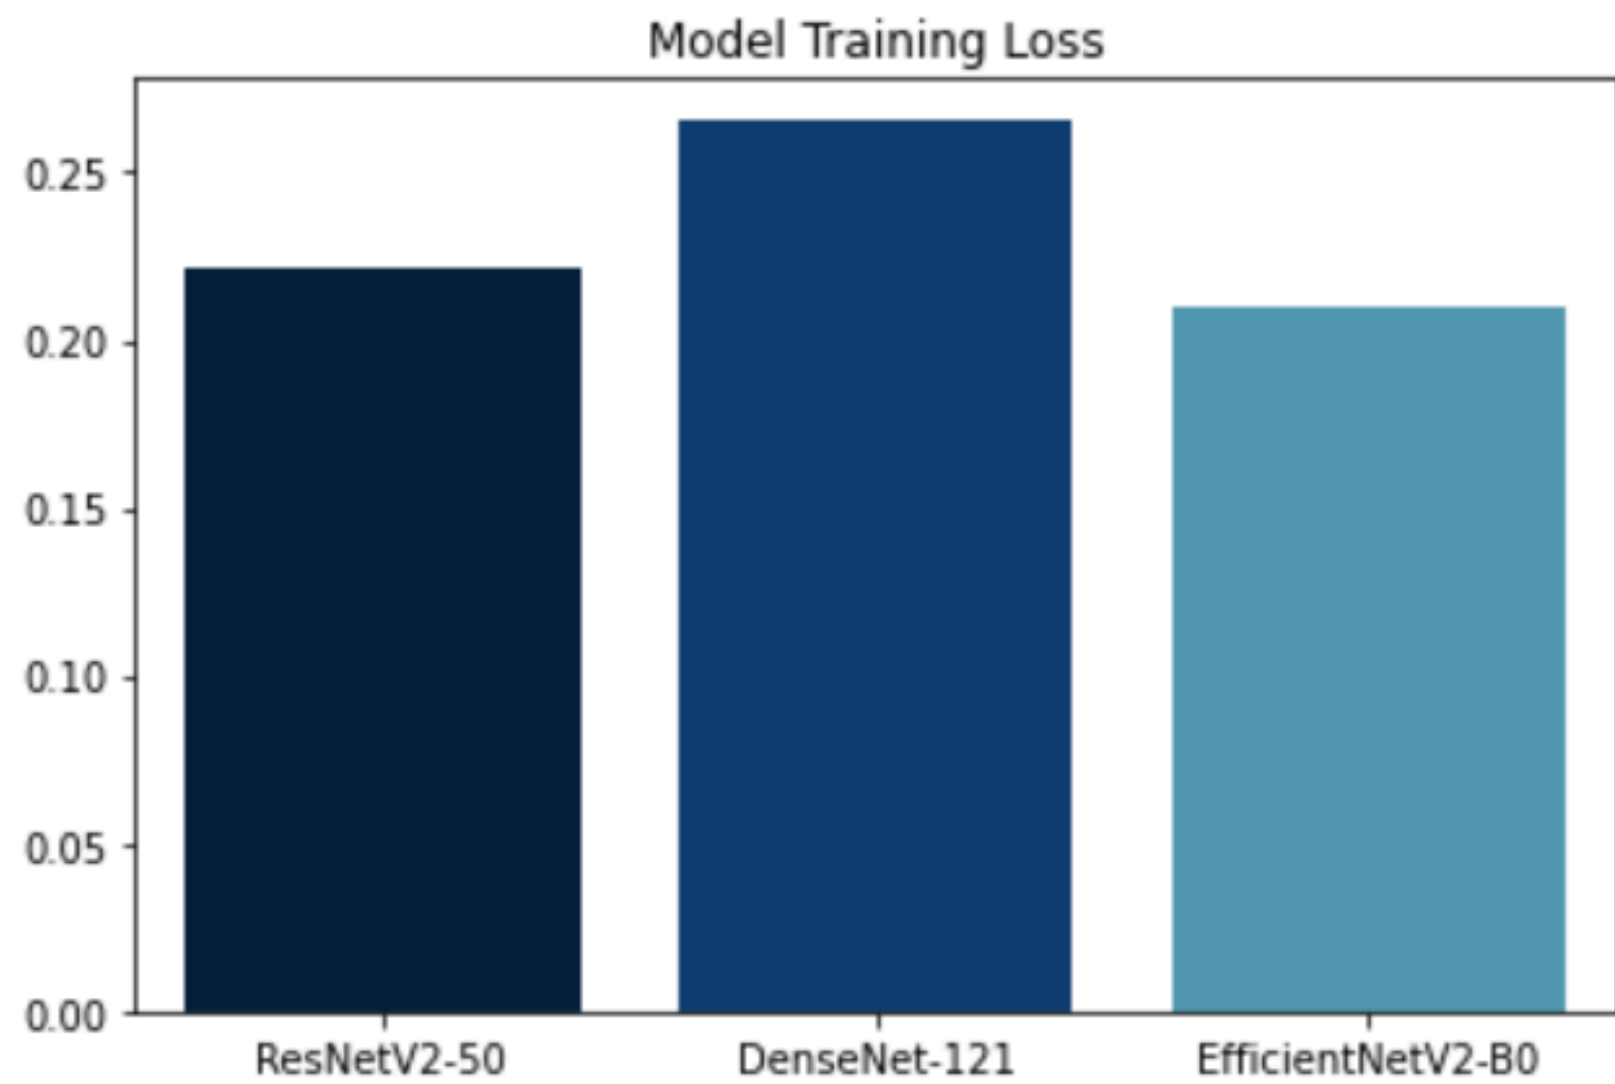


Model Training AUC scores



Model Validation AUC scores





# CHOOSING THE BEST MODEL

**EfficientNetV2-B0**

From the previous slides, we can see that EfficientNetV2-B0 scored the highest in terms of recall, AUC, and loss, so in the next steps, we are going to tune this pretrained model.





# IMPROVING OUR MODEL

## Steps

### ■ Regularization

I added batch normalization and dropout to get a more generalized model.

### ■ Hyperparameter tuning



# Keras

For this tuning process, I used the Hyperband tuner from Keras library, which acts to optimize the search for the right hyperparameters.

# HYPERPARAMETER TUNING RESULTS

```
Trial summary
Hyperparameters:
units: 384
dense_activation: relu
learning_rate: 0.000382870470449823
tuner/epochs: 2
tuner/initial_epoch: 0
tuner/bracket: 3
tuner/round: 0
Score: 0.9448669254779816
```

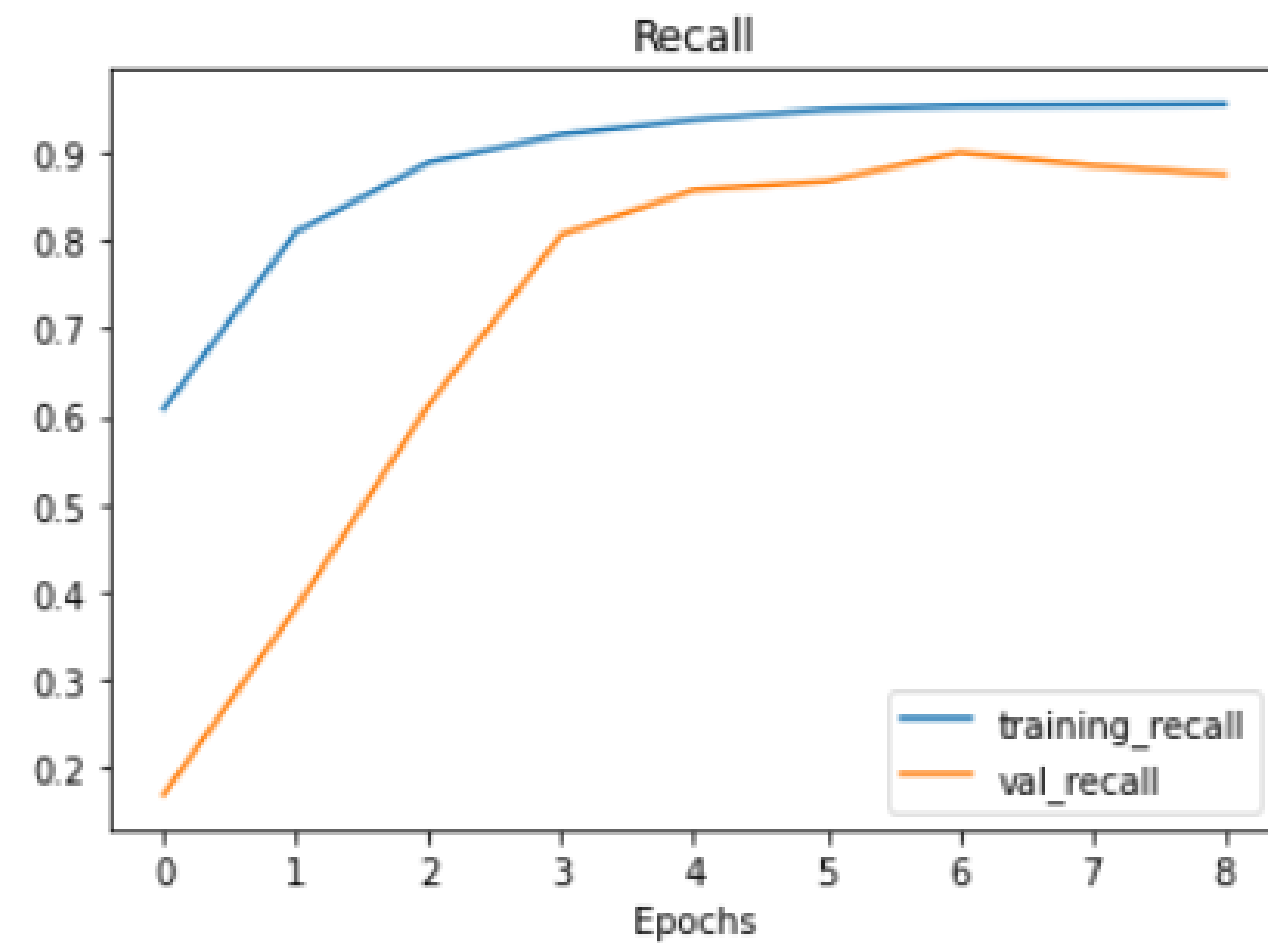

Top #1

```
Trial summary
Hyperparameters:
units: 288
dense_activation: relu
learning_rate: 9.824542588158123e-05
tuner/epochs: 2
tuner/initial_epoch: 0
tuner/bracket: 3
tuner/round: 0
Score: 0.9442332088947296
```

Top #2

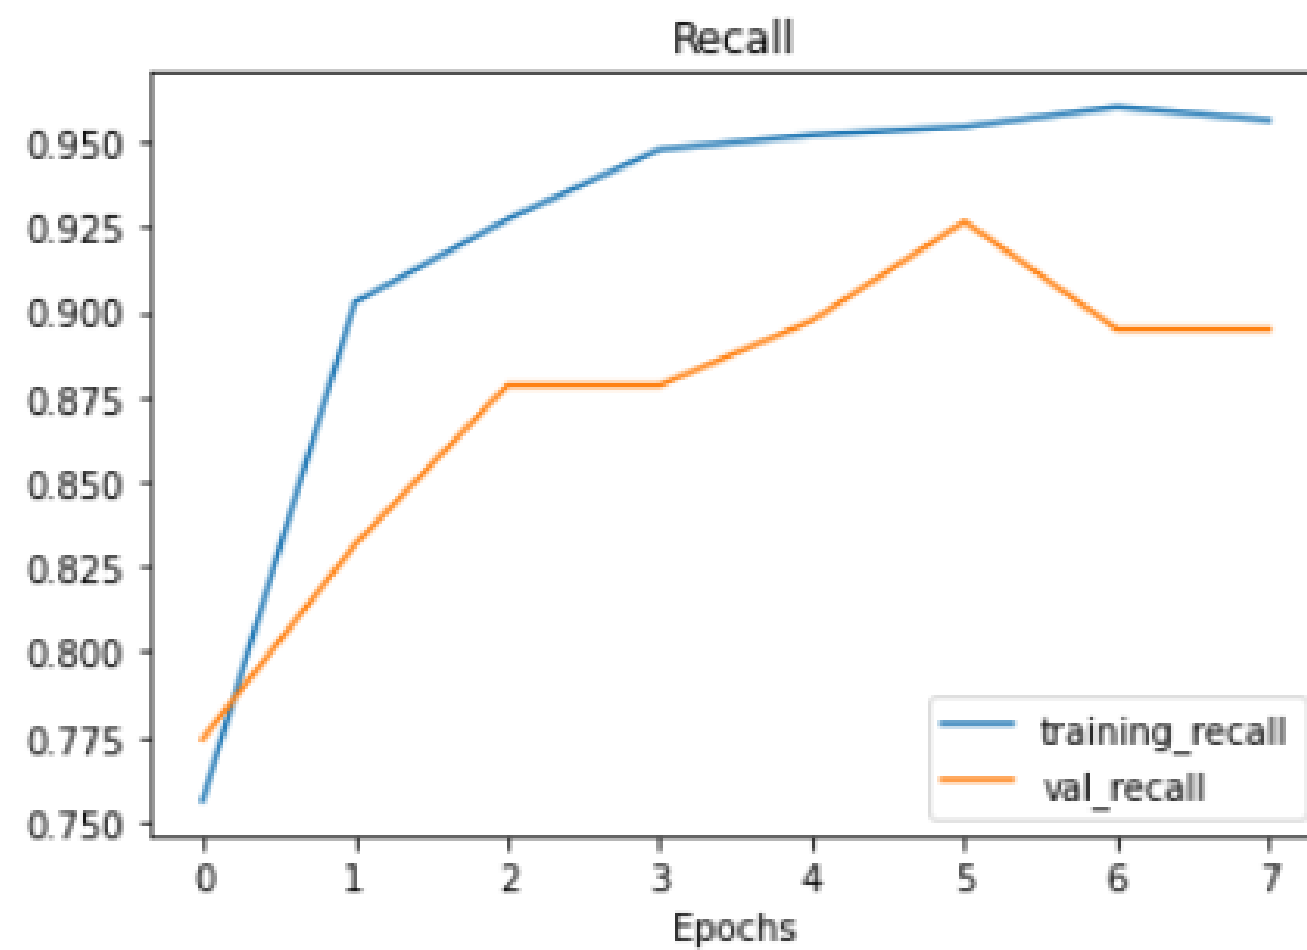

```
Trial summary
Hyperparameters:
units: 64
dense_activation: relu
learning_rate: 2.807830903779318e-05
tuner/epochs: 2
tuner/initial_epoch: 0
tuner/bracket: 3
tuner/round: 0
Score: 0.9391635060310364
```

Top #3



For my tuned model, I decided to take the average of parameters from my initial model and the hyperparameter tuning results. Here are the learning curves of this model.





For comparison, here are the learning curves for my model prior to performing hyperparameter tuning.



# CONFUSION MATRIX



|   | 0   | 1   |
|---|-----|-----|
| 0 | 176 | 58  |
| 1 | 38  | 352 |

## Baseline model

We got 352 True Positives and 38 False Negatives.

|   | 0   | 1   |
|---|-----|-----|
| 0 | 166 | 68  |
| 1 | 6   | 384 |

## Tuned model

We got 384 True Positives and 6 False Negatives

# CLASSIFICATION REPORT

Baseline model

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.82      | 0.75   | 0.79     | 234     |
| 1            | 0.86      | 0.90   | 0.88     | 390     |
| accuracy     |           |        | 0.85     | 624     |
| macro avg    | 0.84      | 0.83   | 0.83     | 624     |
| weighted avg | 0.84      | 0.85   | 0.84     | 624     |

# CLASSIFICATION REPORT

Tuned model

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.97      | 0.71   | 0.82     | 234     |
| 1            | 0.85      | 0.98   | 0.91     | 390     |
| accuracy     |           |        | 0.88     | 624     |
| macro avg    | 0.91      | 0.85   | 0.86     | 624     |
| weighted avg | 0.89      | 0.88   | 0.88     | 624     |



# FINAL REMARKS

Model of choice

Since my tuned model performed the best, I chose that model to be the final one.





# CONCLUSION AND RECOMMEN- DATION





## Conclusion and Recommendation

# CONCLUSION

After testing several models, I concluded that EfficientNetV2-B0 performed the best and achieved a weighted recall score of ~88% and an AUC score of ~94.7% on unseen data, which is an improvement from the baseline model. Our AUC score improved by ~5%, with an improvement of weighted average recall score of ~3%. There was also an improvement of about 30 less cases of false negatives and 30 more cases of true positives!





Conclusion and Recommendation

# RECOMMEN- DATION

The work I have done is far from perfect; there are still many improvements that can be made in the future, such as:

- Spend more time on the hyperparameter tuning process
- Using the ensemble method to create one optimal model





**THANK YOU**